

# Core Algorithm

## 1. Purpose and Context

We need a function that, given a set of anonymized cruise-control samples for a particular road segment plus the official speed-limit rules, produces a single, safe, "best" cruise speed. This chosen speed will drive our simulation and, eventually, real vehicles in adaptive cruise mode.

## 2. Inputs and Outputs

Input	Description
<b>Samples</b> <CruiseSample[]>	A list of records { speedKmh, timestamp }, each one anonymized and stripped of PII
<b>Rules</b> <SegmentRules>	Contains mapSpeedLimitKmh and an optional list of timeDependentLimits (e.g. school-zone hours)
<b>Now</b> <Date> (optional)	The current time, used to evaluate time-dependent limits and filter out stale data

  

Output	Description
<b>Profile</b> <CruiseProfile>	{ segmentId, chosenSpeedKmh, reason } — the final speed recommendation and a human-readable note

## 3. High-Level Steps

### 1. Filter Out Bad Data

- **Age filter:** Discard any sample older than a configurable window (e.g. 30 days).
- **Speed-limit outliers:** Drop any sample exceeding 120 % of the map's speed limit (GPS glitches or driver error).

### 2. Determine the Active Speed Limit

- Start with the official mapSpeedLimitKmh.

- For each defined `timeDependentLimits` window, if the current time falls inside it, take the minimum of that window's limit and the map limit.

### 3. Select the Best Cruise Speed

- **No valid samples:**
  - Fallback immediately to the active speed limit.
  - Reason: "no valid samples."
- **Few samples (< 8):**
  - Take the **most recent** sample's speed (drivers implicitly set what felt best last).
  - Cap it at the active limit.
  - Reason: "only N samples, using most recent."
- **Many samples (≥ 8):**
  - Compute the **median** of all sample speeds (robust to remaining outliers).
  - Cap the median at the active limit.
  - Reason: "median of N samples = X km/h."

### 4. Emit the Cruise Profile

- Package `{ segmentId, chosenSpeedKmh, reason }` and return it.
- 

## 4. Edge Case Considerations

- **All samples too old or too fast:**

Triggers the "no valid samples" branch.

- **Time-window wrap-around** (e.g. night-time restrictions from 22:00 to 06:00):

Ensure your logic correctly handles `startHour > endHour`.

- **Sparse data regions:**

Use a low threshold (e.g. 1 sample) so we can still offer a useful speed, even if it's just the last known driver's choice.

- **Rapidly changing speed limits:**

If a road's posted limit changes (construction, dynamic signage), freshly ingested samples might conflict. You may want to tag samples with the map version they came from and invalidate old-limit samples when rules update.

- **Uniform driver behavior:**

If all drivers always set exactly the limit, median and recent-sample branches converge. That's fine—the algorithm gracefully handles “everyone drives at the limit.”

---

## 5. GDPR & Privacy Compliance

- **Anonymization:** Samples contain **no user identifiers**. Data is aggregated solely by segment.
  - **Aggregation:** We never store or surface individual driver histories—only the chosen profile per segment.
  - **Data retention:** Enforce a rolling window (e.g. 30 days) so old data is purged automatically.
  - **Transparency:** The `reason` field lets us explain how we arrived at each speed, aiding audit and compliance.
- 

## 6. Extensibility & Future Enhancements

- **Weighted sampling:** Add recency weights so newer samples slightly influence the median more.
- **Clustering by context:** Group samples by weather or traffic conditions (if that metadata becomes available) and choose different profiles accordingly.
- **Confidence scoring:** Alongside `chosenSpeedKmh`, compute a confidence metric (e.g. based on sample count or variance) to flag segments that need more data.
- **Dynamic fallback:** For segments with zero samples but low confidence (e.g. rare roads), you could use neighboring segment profiles to estimate a safe speed.